

# GAGAN SAINI

[sainigagan163@gmail.com](mailto:sainigagan163@gmail.com)

+91 9911774256

Delhi - NCR

[LinkedIn](#)

[Portfolio](#)

[GitHub](#)

## PROFESSIONAL SUMMARY

Data Scientist with a track record of building ML systems that reduce cost, automate complexity, and surface actionable insights from large-scale operational data - including a clustering model cutting client cloud spend by 25% and an anomaly detection system reducing incident response time by 30%. Currently pursuing MSc in Data Science at the University of Aberdeen (graduating June 2026), with 2.5 years of commercial experience in Python, SQL, and cloud-based data engineering across AWS, Azure, and GCP. Holds five certifications including Azure Data Scientist Associate.

## TECHNICAL SKILLS

Python, SQL, Machine Learning, Statistical Modelling, Scikit-learn, TensorFlow, XGBoost, Pandas, NumPy, Matplotlib, Exploratory Data Analysis, Feature Engineering, Predictive Modelling, ETL Pipeline Development, Azure, AWS, GCP, Git, Version Control

## EDUCATION

MSc Data Science | University of Aberdeen, UK | Jan 2025 - Jun 2026 (Expected)

Bachelor of Technology | Dr. A.P.J. Abdul Kalam Technical University | Aug 2017 - Sep 2021

## EXPERIENCE

### Data Engineer I

07/2022 - 12/2024

#### CloudEQ Software India Pvt. Ltd.

- Developed a machine learning-powered resource clustering model in Python using Scikit-learn, TensorFlow, and DBScan to perform exploratory analysis and identify utilisation patterns across 500+ cloud assets, surfacing cost-reduction signals that contributed to a 25% reduction in client cloud expenditure.
- Built automated ETL pipelines using Python and SQL to collect, clean, and transform high-volume infrastructure telemetry from AWS, Azure, and GCP APIs into structured analytical datasets - reducing manual data preparation overhead by 40% and enabling consistent downstream statistical modelling.
- Engineered a Python-based anomaly detection system applying threshold-based statistical rules to live operational metrics, identifying performance degradation patterns across 500+ monitored assets before escalation and reducing incident response time by 30%.
- Applied unsupervised clustering techniques to multi-source cloud utilisation data to segment resource behaviour, producing structured analytical outputs that informed infrastructure optimisation decisions across client environments.
- Developed and maintained data pipelines using Microsoft Fabric Data Factory, automating structured data ingestion workflows and contributing to scalable, reproducible analytical infrastructure.

## PROJECT HIGHLIGHTS

- Built a scalable synthetic fraud detection dataset generator in Python producing 100K+ rows across a 20-column schema, modelling four distinct fraud patterns using a VocabularyFactory class and DatasetConfig dataclass - demonstrating end-to-end ML data pipeline design and reproducible analytical methodology.
- Developed an E-sports Match Winner Predictor using XGBoost on 50K+ League of Legends matches, achieving 94.6% accuracy and 0.986 ROC-AUC through feature engineering on first-15-minute in-game telemetry.

## CERTIFICATIONS

Microsoft Certified: Azure AI Engineer Associate | Microsoft Certified: Azure Data Scientist Associate |  
AWS Certified DevOps Engineer - Professional | AWS Certified Developer - Associate |  
AWS Certified Cloud Practitioner | Full Stack Python Diploma - FirstBit Solutions